



SOCKET PROGRAMING

IN LINUX - C

Courtesy of Roni Levin, Yoad Zur & Chen Griner

Introduction

Why learn about sockets?

- The basics of app-to-app communication over the internet.

What is the goal of this presentation?

- **Introduce** you with some concepts of socket programming. “ How to work with network sockets? “

This is a very wide subject

- This is not a complete guide! You are expected to explore further by yourself.

Today's Presentation:


- 1 • High-level Overview, clients and servers
- 2 • Familiarize yourself with socket API
- 3 • A simple socket connection scheme
- 5 • Socket management
- 6 • Lab Session



HIGH-LEVEL OVERVIEW: CLIENTS AND SERVERS

Client – Server

- **Clients:**
 - Locates the server
 - Initiate connection
 - Example: you.
- **Server:**
 - Responder.
 - Provides service.
 - Example: Moodle.



Client – *Server*: some key differences

Clients

- ☐ Simple
- ☐ (Usually) sequential
- ☐ Not performance sensitive
- ☐ **Execute on-demand**

Server

- ☐ Complex
- ☐ (Massively) concurrent
- ☐ High performance
- ☐ **Always-on**

Client – *Server*: Similarities

- ❑ Share common protocols
 - Network layer
 - Transport layer
 - Application layer
- ❑ Both rely on APIs for network access

What is an API?

application programming interface (API) is a set of routines, protocols, and tools for building software applications. For example network sockets

What is a network socket?

It is an application's
“mailbox” for
network messages.

Used to pass
messages among
applications on
different computers.

Managed by the
operating system.

Represented as a
“*file descriptor*”.

Implements an
Incoming and
Outgoing queues.

Identified by IP
address, port and
protocol.

Client: general workflow

Create a socket

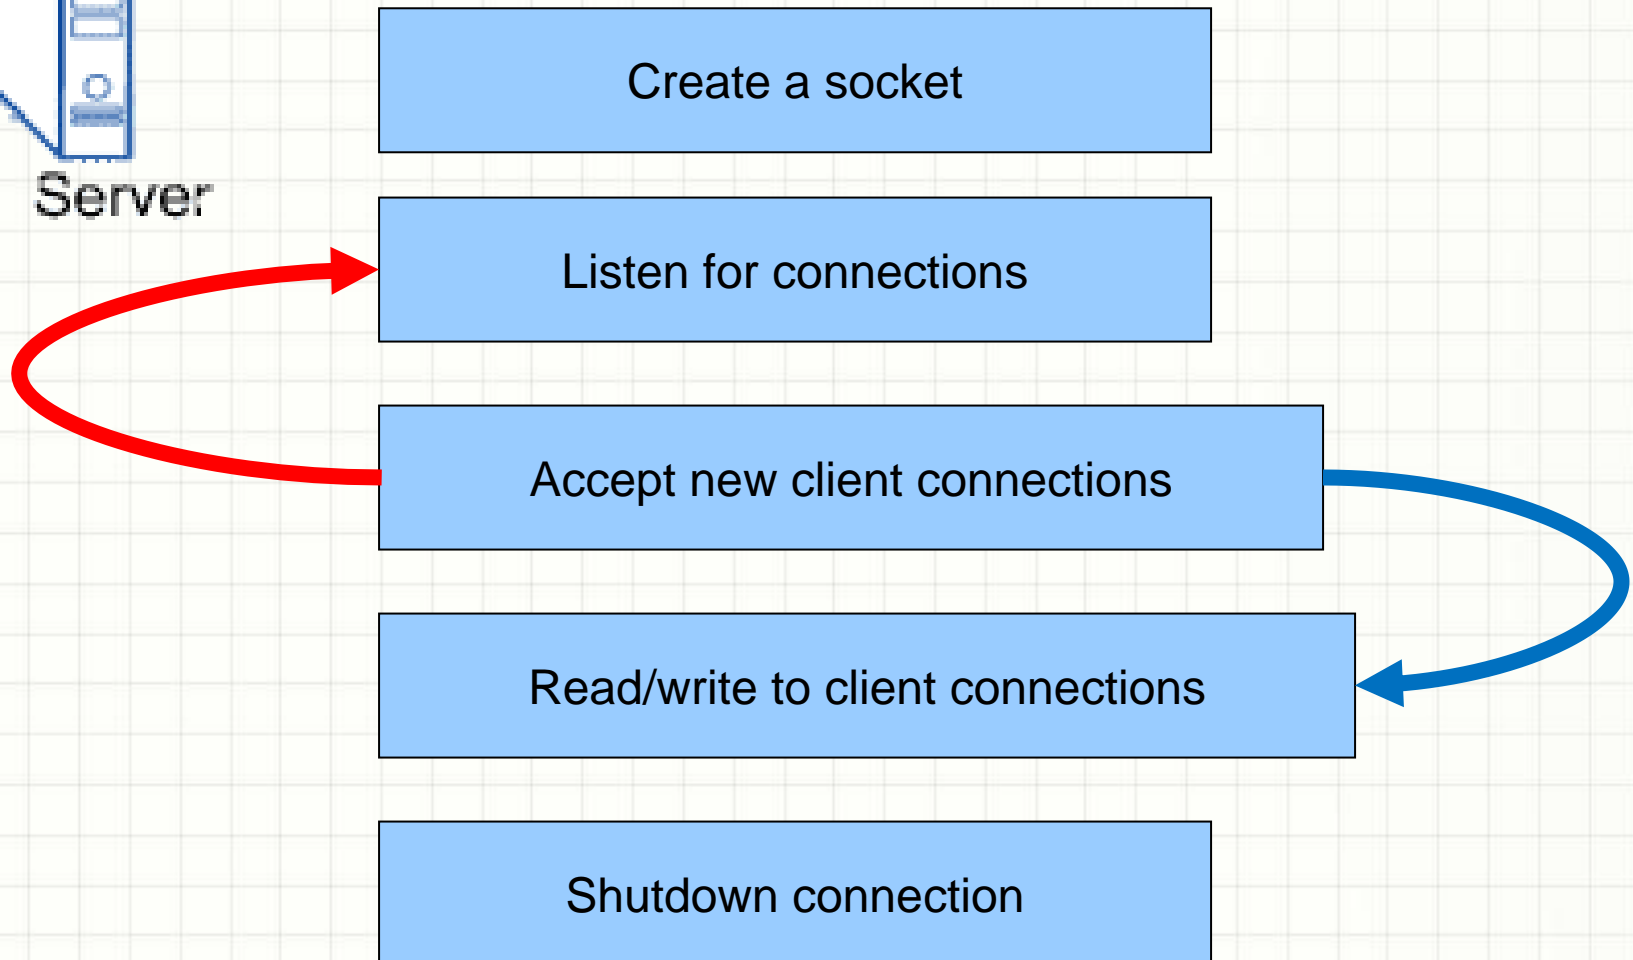

Find/use the server address

Connect to the server

Read/write data

Shutdown connection

Server: general work flow



TCP *Server*:

The “Welcome socket”

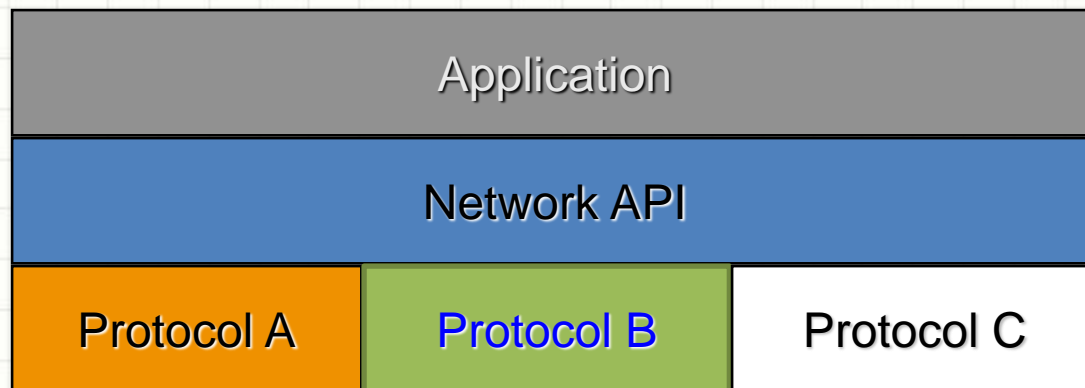
- The client needs a known address and port to connect to.
- A TCP server has at least one special socket we call the “welcome socket”.
- First connection by the client is to this socket.
- New separate socket for each new client.
- The “welcome socket” is permanent.
 - It is not closed when the client leaves.
- The welcome socket has a queue.
 - Can’t have too many clients entering at the same time.



SOCKET API

Introduction

- A **socket API** is an application programming interface (API), that allows applications to control and use network sockets.
- Sockets provide the interface between application and the protocol software.



Berkeley sockets / POSIX sockets

- Used in Unix systems since 1983.
- We will run our programs in Linux.
- An API for Internet sockets.
 - Essentially, you get a set of C functions.
 - Each function has a different role in the operation of sockets.
 - What functions to use?

Where do we start?

Basic example, opening a socket

- Let's first declare two sockets, one for UDP and one for TCP.
 - (1) A socket is a file-descriptor. An **integer**.
- Let's tell the OS we want to make a new socket.
 - (2) we use the `socket()` function and get a TCP socket.
 - (3) same for UDP.
- `SOCK_STREAM` Or `SOCK_DGRAM` determine the socket type.
- Is that it?
 - There are many other parts.

```
1. int myTCPsocket, myUDPsocket;  
2. myTCPsocket = socket(AF_INET, SOCK_STREAM, 0);  
3. myUDPsocket = socket(AF_INET, SOCK_DGRAM, 0);
```

Socket Types

- The stream socket
 - Data should be handled as a part of a stream
 - Maintain a constant connection
 - Uses TCP to set up a reliable connection
 - Use: `SOCK_STREAM`
- The datagram socket.
 - Data is separate packets.
 - Use UDP
 - Unreliable
 - Use: `SOCK_DGRAM`
- Another
 - Raw socket
 - (don't) Use: `SOCK_RAW`

Socket API

- `socket()`: creates a socket of a given domain, type, protocol
- <http://linux.die.net/man/2/socket>

```
1. TCP: socket(AF_INET, SOCK_STREAM, 0)
2. UDP: socket(AF_INET, SOCK_DGRAM, 0)
```

- `bind()`: a socket start with no address, assigns an address to the socket.
 - Remember to fill and cast an appropriate “sockaddr” type struct.
- <http://linux.die.net/man/2/bind>

```
1. TCP server & UDP: bind(lisen_sock, (struct sockaddr *)&listen_addr, sizeof(listen_addr))
```

Ipv4 socket address structure

```
struct sockaddr_in{
    uint8_t      sin_len; /*length of the structure (16)*/
    sa_family_t  sin_family /* AF_INET*/
    in_port_t    sin_port /* 16 bit TCP or UDP port number*/
    struct in_addr sin_addr /* 32 bit Ipv4 address */
    char         sin_zero(8)/* unused*/
}
```

Socket API

- `listen()`: specifies the number of pending concurrent connections that can be queued for a server socket.
 - Makes a “welcome socket”
- <http://linux.die.net/man/2/listen>

```
1. TCP: listen(lisen_sock, SOMAXCONN)
```

Socket API

- `accept()`: server accepts a connection request from a client. Blocking until a connection is received.
- <http://linux.die.net/man/2/accept>

```
1. TCP: accept(lisen_sock, (struct sockaddr*) &client_addr, &client_addr_size)
```

- `inet_addr()`: converts the Internet host address *cp* from IPv4 numbers-and-dots notation into binary data in network byte order
 - Another similar function `gethostbyname()`.
- https://linux.die.net/man/3/inet_addr

```
1. TCP or UDP: in_addr_t inet_addr(const char *cp);
```

- `connect()`: client requests a connection request to a server (call)
- <http://linux.die.net/man/2/connect>

```
1. TCP: connect(server_sock, (struct sockaddr*) &server_addr, sizeof(server_addr))
```

Socket API

- `recv()`, `recvfrom()`: read from a connection
 - Positive returned value marks the number of successfully received bytes
 - negative returned value means an error.
 - 0 means the connection has been closed.
 - On which protocol is this relevant?
- <http://linux.die.net/man/2/recv>

```
1. TCP & UDP : recv(sock, buffer, &buff_len, 0)
```

- `send()`, `sendto()`: write to a connection
 - Positive returned value marks the number of sent bytes.
 - What happens if not all bytes were sent?
 - negative returned value means an error.
- <http://linux.die.net/man/2/send>

```
1. TCP & UDP : send(sock, buffer, &buff_len, 0)
```


Socket API

- `Setsockopt(...)`: assigns options to the socket.
 - Join/leave multicast group.
 - Other options.
 - Use as necessary.
- <https://linux.die.net/man/2/setsockopt>

TCP & UDP :

```
int setsockopt(int socket, int level, int option_name, const void *option_value, socklen_t option_len);
```

Socket API

- `Shutdown(...)`. shut down socket send and receive operations.
 - A socket is full duplex.
 - Stop receiving or stop sending.
 - You may just use `close(..)`
- <https://linux.die.net/man/3/shutdown>

TCP & UDP:

```
int shutdown(int socket, int how);
```

- `Close(...)`: close the socket.
 - What happens when you close a TCP/UDP socket?
- <http://linux.die.net/man/2/close>

```
1. TCP & UDP: close(socket)
```

Socket API

- `htons()`, `htonl()`, `ntohs()`, `ntohl()`: Convert to or from host byte order to network byte order, i.e., big endian \Leftrightarrow little endian.
 - `htons()`: host to network short.
 - `htonl()`: host to network long.
 - `ntohs()`: network to host short.
 - `ntohl()`: network to host long.
- <http://linux.die.net/man/3/htons>

`1. TCP & UDP : htons(1234); // 1234=port number`

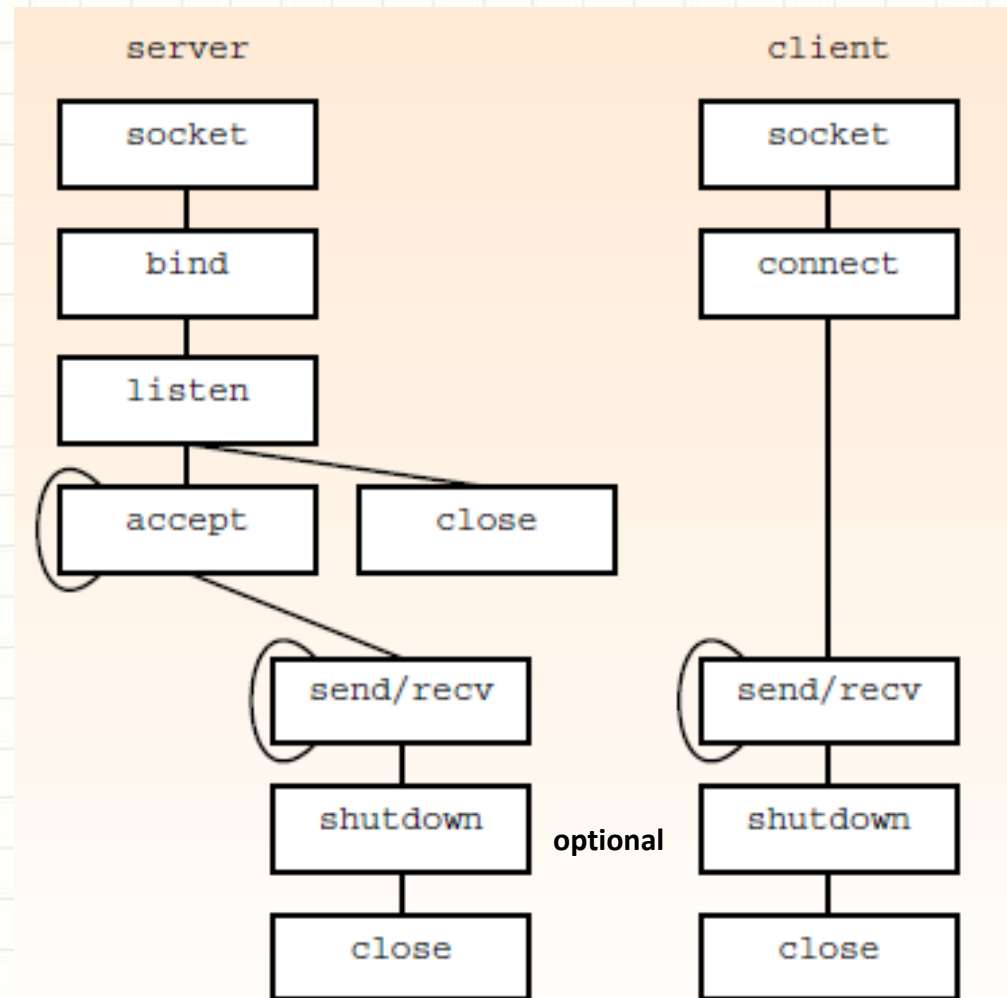
There are many more socket related functions

- Extra: a nice example of client-server program server using socket API:
- <https://www.programminglogic.com/example-of-client-server-program-in-c-using-sockets-and-tcp/>

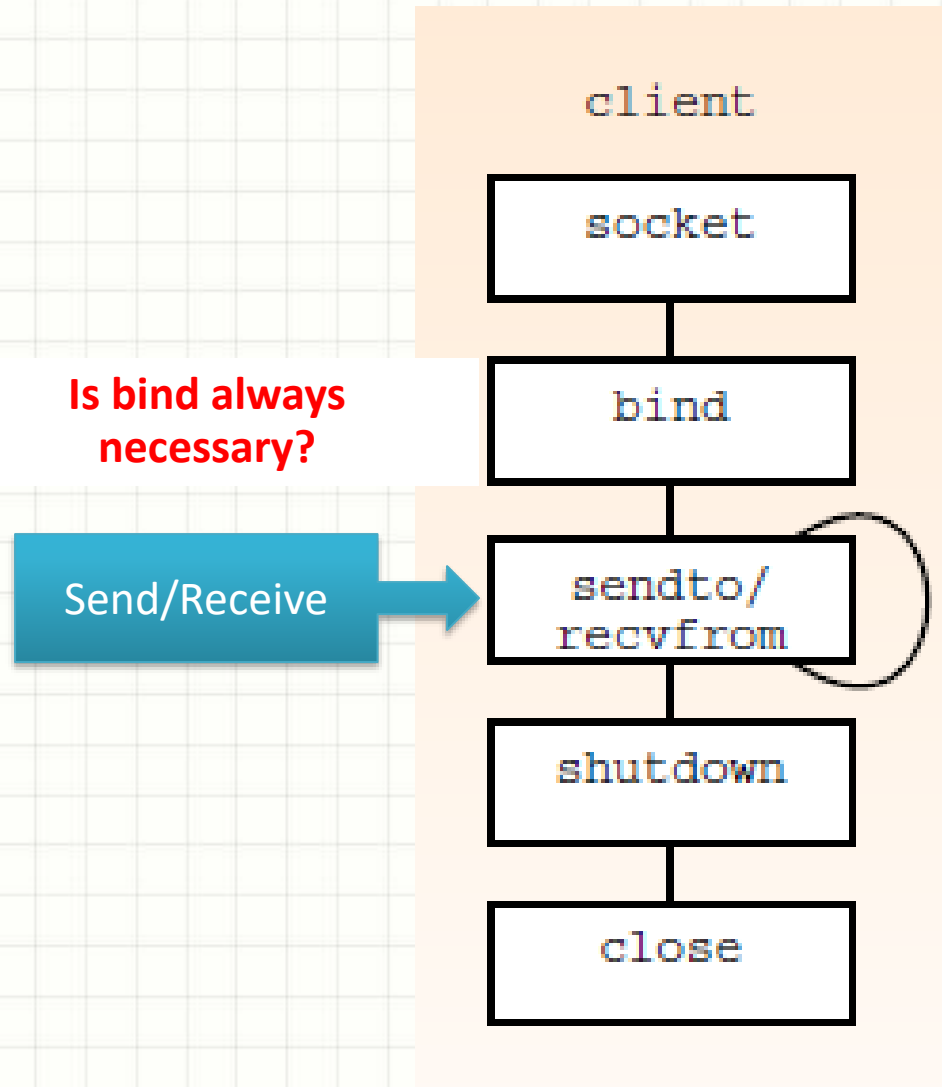


SOCKET CONNECTION SCHEME

TCP



UDP





SOCKET MANAGEMENT

Socket management methods

- Servers need to manage many socket simultaneously.
 - At least one “welcome socket”.
 - Many clients.
 - How?
- Parallel: Threads / Process.
- Serial: `select()` function

Serial Socket management: Select()

- **Select()**: waits for sockets(or any FD) to change status
 - To use **select()** first define a set of FDs (fd_set).
 - readfds :a set of *read* descriptors to monitor .
 - writefds :a set of *write* descriptors to monitor
- **Select()** start blocking until an event occurs.
 - An event could be an incoming message, a timeout and more.
- <http://linux.die.net/man/2/select>
- Read more about select:
<http://www.lowtek.com/sockets/select.html>

```
TCP & UDP : int select(int nfd, fd_set *readfds, fd_set  
*writefds, fd_set *exceptfds, struct timeval *timeout);
```

```
void FD_CLR(int fd, fd_set *set);  
int FD_ISSET(int fd, fd_set *set);  
void FD_SET(int fd, fd_set *set);  
void FD_ZERO(fd_set *set);
```

Select(); Example

The following is a single use case for select().

We listen for user input until a time out has occurred or the user has entered some input.

Note that this example is incomplete.

TCP & UDP & files

```
1. fd_set fdset; // set of file descriptors
2. FD_ZERO(& fdset); // clear the set
3. FD_SET(fileno(stdin), & fdset); // set stdin as a file descriptor in the set.
4. timeout = set_timeout(60 sec) // set the timeout for select, not a "real" function.
5. inputfd = select(FD_SETSIZE, & fdset, NULL, NULL, & timeout) // block until timeout, or user input
6.     if(inputfd == 0)
7.         { puts("Oops, Time Out!!!") } // not a "real" function
8.     else if(FD_ISSET(fileno(stdin), & fdset)
9.         { user_handler() } // not a "real" function
```

Select(); Example

Some questions?

- Do we have to use a timeout?
 - No.
- Can we listen to multiple sockets using the same FD_set?
 - Yes.
- Can we have more than one file descriptor “jumping” in the same time.
 - Yes.
- Will this select() work multiple times?
 - Not without some further work.

To understand how to work with select() you will have to do some further research on your own.

Use example, select()


- Assume the following scenario:
 - A client program sending a random message once in some random interval of time, repeatedly.
 - Each client has a unique ID.
 - A server program. The clients connect to the server, every time that the server receives a message, it prints “Client with ID X has sent a message”.
 - Several clients connect to the server at once.
- How can the server be implemented using select() ?



SOCKET PROGRAMMING

Socket Programing

- In the first part, we will learn the fundamentals of socket programming.
- We will prepare 5 different short programs.
- Two for TCP.
 - A TCP receiver.
 - A TCP sender.
- Two for UDP
 - A UDP sender.
 - A UDP receiver capable of listening to a UDP stream.
- A program that can listen to several UDP sockets at once.



Socket Programming

- Each sender program sends a number of messages with data from a file.
- Each receiver program receives this data and prints it.
- Each program will run on the lab's PCs, connected by a GNS3 topology from previous labs.
- We will use a new virtual machine with a GUI, PC5 – either write your programs in PC5 or use it as a middleman to a cloud.

Connect it to a GNS3 topology and transfer your programs to the 4 lab PCs (via FTP or SCP)

Thanks to...(and other helpful links)

- Wikipedia
- http://www.cs.northwestern.edu/~agupta/cs340/sockets/sockets_intro.ppt
- <http://parsys.eecs.uic.edu/~solworth/sockets.pdf>
- <https://www.cs.cmu.edu/~srini/15-441/S10/lectures/r01-sockets.pdf>